

Data Management in Large-Scale Distributed Systems

NoSQL Databases

Thomas Ropars

`thomas.ropars@univ-grenoble-alpes.fr`

`http://tropars.github.io/`

2019

References

- The lecture notes of V. Leroy
- The lecture notes of F. Zanen Boito
- Designing Data-Intensive Applications by Martin Kleppmann
 - ▶ Chapter 7

In this lecture

- Motivations for NoSQL databases
- ACID properties and CAP Theorem
- A landscape of NoSQL databases

Agenda

Introduction

Why NoSQL?

Transactions, ACID properties and CAP theorem

Data models

NoSQL databases design and implementation

Common patterns of data accesses

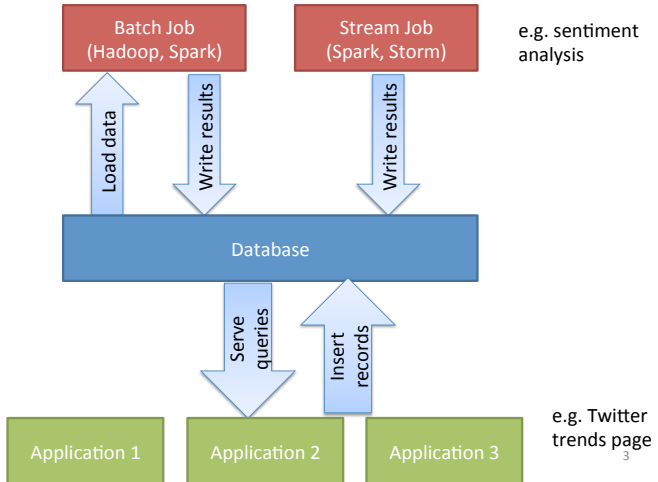
Large-scale data processing

- Batch processing: Hadoop, Spark, etc.
- Perform some computation/transformation over a full dataset
- Process all data

Selective query

- Access a specific part of the dataset
- Manipulate only data needed (1 record among millions)
- Main purpose of a database system

Processing / Database Link



Different types of databases

- So far we used HDFS



- A file system can be seen as a very basic database
- Directories / files to organize data
- Very simple queries (file system path)
- Very good scalability, fault tolerance ...

- Other end of the spectrum: Relational Databases

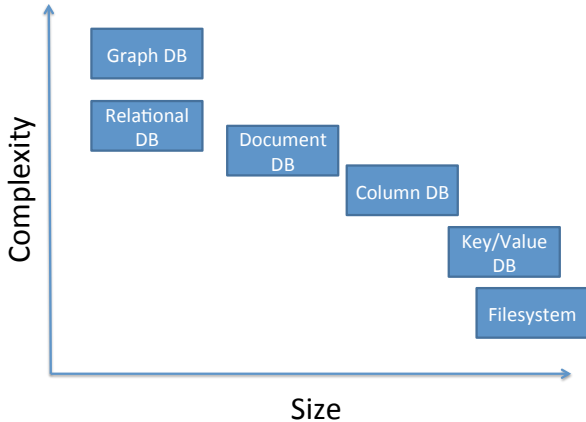
- SQL query language, very expressive
- Limited scalability (generally 1 server)

















PostgreSQL



Size / Complexity



The NoSQL Jungle

Document Database	Graph Databases
  	 
Wide Column Stores	Key-Value Databases
   	    

@cloudbt <http://www.anyannava.com>

Agenda

Introduction

Why NoSQL?

Transactions, ACID properties and CAP theorem

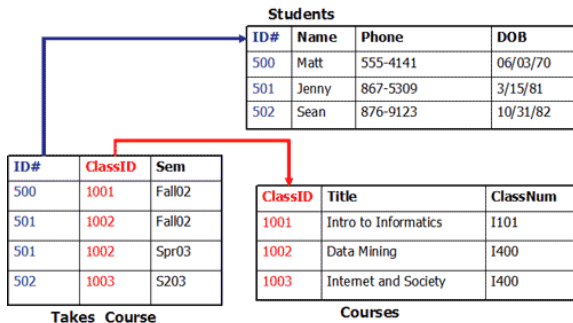
Data models

NoSQL databases design and implementation

Relational databases

SQL

- Born in the 70's – Still heavily used
- Data is organized into relations (in SQL: tables)
- Each relation is an unordered collection of tuples (rows)



About SQL

Advantages

- Separate the data from the code
 - ▶ High-level language
 - ▶ Space for optimization strategies
- Powerful query language
 - ▶ Clean semantics
 - ▶ Operations on sets
- Support for transactions

Motivations for alternative models

see <https://blog.couchbase.com/nosql-adoption-survey-surprises/>

Some limitations of relational databases

- Performance and scalability
 - ▶ Difficult to partition the data (in general run on a single server)
 - ▶ Need to scale up to improve performance
- Lack of flexibility
 - ▶ Will to easily change the schema
 - ▶ Need to express different relations
 - ▶ Not all data are well structured
- Few open source solutions
- Mismatch between the relational model and object-oriented programming model

Illustration of the object-relational mismatch

Figure by M. Kleppmann

<http://www.linkedin.com/in/williamhgates>



Bill Gates
Greater Seattle Area | Philanthropy

Summary
Co-chair of the Bill & Melinda Gates Foundation. Chairman, Microsoft Corporation. Voracious reader. Avid traveler. Active blogger.

Experience
Co-chair • Bill & Melinda Gates Foundation
2000 – Present
Co-founder, Chairman • Microsoft
1975 – Present

Education
Harvard University
1973 – 1975
Lakeside School, Seattle

Contact Info
Blog: thegatesnotes.com
Twitter: @BillGates

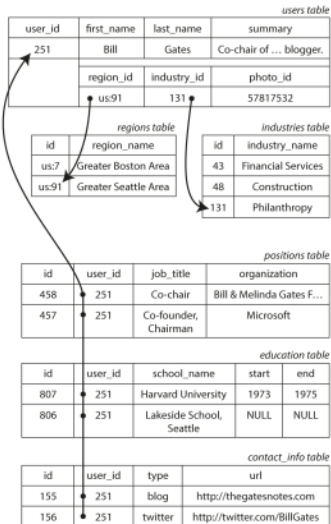


Figure: A CV in a relation database

Illustration of the object-relational mismatch

Figure by M. Kleppmann

```
{
  "user_id": 251,
  "first_name": "Bill",
  "last_name": "Gates",
  "summary": "Co-chair of the Bill & Melinda Gates; Active blogger.",
  "region_id": "us:91",
  "industry_id": 131,
  "photo_url": "/p/7/000/253/05b/308dd6e.jpg",
  "positions": [
    {"job_title": "Co-chair", "organization": "Bill & Melinda Gates Foundation"},
    {"job_title": "Co-founder, Chairman", "organization": "Microsoft"}
  ],
  "education": [
    {"school_name": "Harvard University", "start": 1973, "end": 1975},
    {"school_name": "Lakeside School, Seattle", "start": null, "end": null}
  ],
  "contact_info": {
    "blog": "http://thegatesnotes.com",
    "twitter": "http://twitter.com/BillGates"
  }
}
```

Figure: A CV in a JSON document

About NoSQL

What is NoSQL?

- A hashtag
 - ▶ NoSQL approaches were existing before the name became famous
- No SQL
- New SQL
- Not only SQL
 - ▶ Relational databases will continue to exist alongside non-relational datastores

About NoSQL

A variety of NoSQL solutions

- Key-Value (KV) stores
- Wide column stores (Column family stores)
- Document databases
- Graph databases

Difference with relational databases

There are several ways in which they differ from relational databases:

- Properties
- Data models
- Underlying architecture

Agenda

Introduction

Why NoSQL?

Transactions, ACID properties and CAP theorem

Data models

NoSQL databases design and implementation

About transactions

The concept of transaction

- Groups several read and write operations into a logical unit
- A group of reads and writes are executed as one operation:
 - ▶ The entire transaction succeeds (commit)
 - ▶ or the entire transaction fails (abort, rollback)
- If a transaction fails, the application can safely retry

About transactions

The concept of transaction

- Groups several read and write operations into a logical unit
- A group of reads and writes are executed as one operation:
 - ▶ The entire transaction succeeds (commit)
 - ▶ or the entire transaction fails (abort, rollback)
- If a transaction fails, the application can safely retry

Why do we need transactions?

- Crashes may occur at any time
 - ▶ On the database side
 - ▶ On the application side
 - ▶ The network might not be reliable
- Several clients may write to the database at the same time

ACID

ACID describes the set of safety guarantees provided by transactions

- Atomicity
- Consistency
- Isolation
- Durability

Having such properties make the life of developers easy, but:

- ACID properties are not the same in all databases
 - ▶ It is not even the same in all SQL databases
- NoSQL solutions tend to provide weaker safety guarantees
 - ▶ To have better performance, scalability, etc.

ACID: Atomicity

Description

- A transactions succeeds completely or fails completely
 - ▶ If a single operation in a transaction fails, the whole transaction should fail
 - ▶ If a transaction fails, the database is left unchanged
- It should be able to deal with any faults *in the middle* of a transaction
- If a transaction fails, a client can safely retry

In the NoSQL context:

- Atomicity is still ensured

ACID: Consistency

Description

- Ensures that the transaction brings the database from a valid state to another valid state
 - ▶ Example: Credits and debits over all accounts must always be balanced
- It is a property of the application, not of the database
 - ▶ The application cannot enforce application-specific invariants
 - ▶ The database can check some specific invariants
 - A foreign key must be valid

In the NoSQL context:

- Consistency is (often) not discussed

ACID: Durability

Description

- Ensures that once a transaction has committed successfully, data will not be lost
 - ▶ Even if a server crashes (flush to a storage device, replication)

In the NoSQL context:

- Durability is also ensured

ACID: Isolation

Description

- Concurrently executed transactions are isolated from each other
 - ▶ We need to deal with concurrent transactions that access the same data
- **Serializability**
 - ▶ High level of isolation where each transaction executes as if it was the only transaction applied on the database
 - As if the transactions are applied *serially*, one after the other
 - ▶ Many SQL solutions provide a lower level of isolation

In the NoSQL context:

- **What about the CAP theorem?**

The CAP theorem

3 properties of databases

- Consistency
 - ▶ What guarantees do we have on the value returned by a read operation?
 - ▶ It strongly relates to Isolation in ACID (and not to consistency)
- Availability
 - ▶ The system should always accept updates
- Partition tolerance
 - ▶ The system should be able to deal with a partitioning of the network

Comments on CAP theorem

- Was introduced by E. Brewer in its lectures (beginning of years 2000)
- Goal: discussing trade-offs in database design

What does the CAP theorem says?


The theorem

It is impossible to have a system that provides Consistency, Availability, and partition tolerance.

How it should be understood:

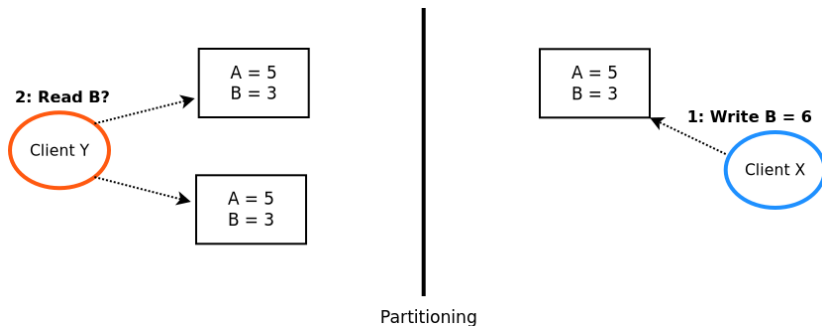
- Partitions are unavoidable
 - ▶ It is a fault, we have no control on it
- We need to choose between availability and consistency
 - ▶ In the CAP theorem:
 - Consistency is meant as *linearizability* (the strongest consistency guarantee)
 - Availability is meant as *total availability*
 - ▶ In practice, different trade-offs can be provided

The intuition behind CAP



- Let inconsistencies occur? (No C)
- Stop executing transactions? (No A)

The intuition behind CAP



- Let inconsistencies occur? (No C)
- Stop executing transactions? (No A)

Note that in a centralized system (non-partitioned relational database), no need for Partition tolerance

- We can have Consistency and Availability

The impact of CAP on ACID for NoSQL

source: E. Brewer

The main consequence

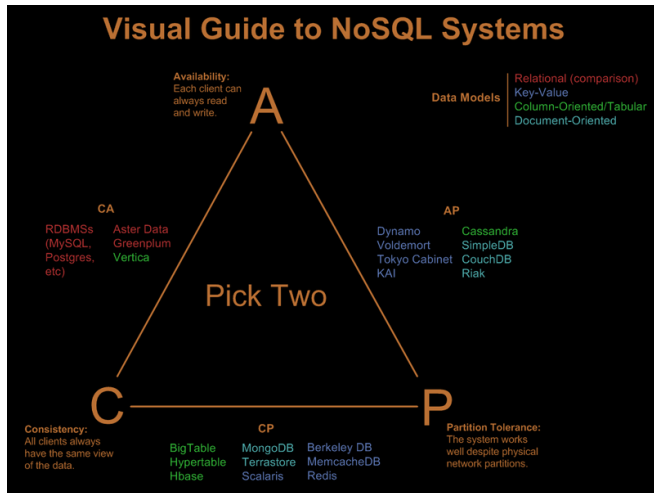
- No NoSQL database with strong Isolation

Discussion about other ACID properties

- Atomicity
 - ▶ Each side should ensure atomicity
- Durability
 - ▶ Should never be compromised

A vision of the NoSQL landscape

Source: <https://blog.nahurst.com/visual-guide-to-nosql-systems>



To be read with care:

- Solutions often provide a trade-off between CP and AP
- A single solution may often a different trade-off depending on how is is configured.
- **We don't pick two !**

Agenda

Introduction

Why NoSQL?

Transactions, ACID properties and CAP theorem

Data models

NoSQL databases design and implementation

Agenda

Introduction

Why NoSQL?

Transactions, ACID properties and CAP theorem

Data models

NoSQL databases design and implementation

Additional references

Mandatory reading

- *Bigtable: A Distributed Storage System for Structured Data.*, F. Chang et al., OSDI, 2006.
- *Cassandra: a decentralized structured storage system* ., A. Lakshman et al., SIGOPS OS review, 2010.

Suggested reading

- <http://martin.kleppmann.com/2015/05/11/please-stop-calling-databases-cp-or-ap.html>, M. Kleppmann, 2015.
- <https://jvns.ca/blog/2016/11/19/a-critique-of-the-cap-theorem/>, J. Evans, 2016.