

# Data Management in Large-Scale Distributed Systems

NoSQL Databases

Thomas Ropars

`thomas.ropars@univ-grenoble-alpes.fr`

`http://tropars.github.io/`

2019

# References

- The lecture notes of V. Leroy
- The lecture notes of F. Zanon Boito
- Designing Data-Intensive Applications by Martin Kleppmann
  - ▶ Chapters 2 and 7

## In this lecture

- Motivations for NoSQL databases
- ACID properties and CAP Theorem
- A landscape of NoSQL databases

# Agenda

Introduction

Why NoSQL?

Transactions, ACID properties and CAP theorem

Data models

NoSQL databases design and implementation

# Common patterns of data accesses

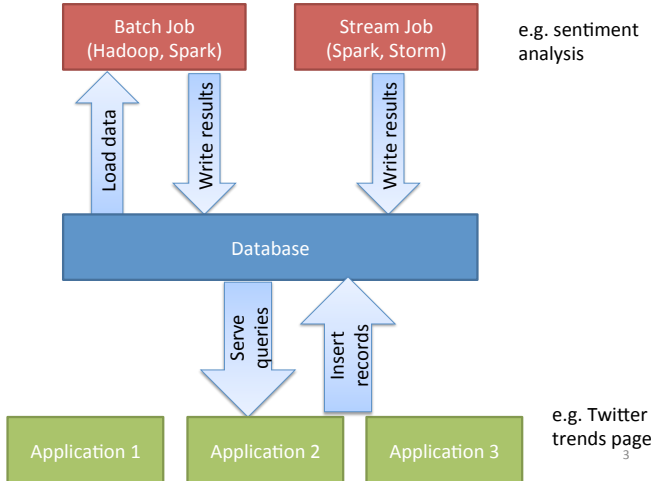
## Large-scale data processing

- Batch processing: Hadoop, Spark, etc.
- Perform some computation/transformation over a full dataset
- Process all data

## Selective query

- Access a specific part of the dataset
- Manipulate only data needed (1 record among millions)
- Main purpose of a database system

# Processing / Database Link



# Different types of databases

- So far we used HDFS



- A file system can be seen as a very basic database
- Directories / files to organize data
- Very simple queries (file system path)
- Very good scalability, fault tolerance ...

- Other end of the spectrum: Relational Databases

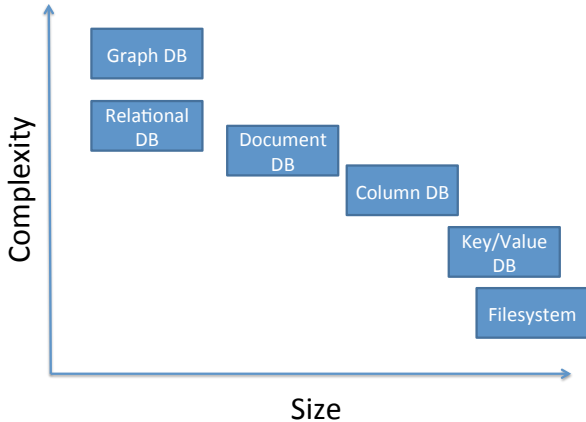


PostgreSQL

- SQL query language, very expressive
- Limited scalability (generally 1 server)

















# Size / Complexity





# The NoSQL Jungle

Document Database	Graph Databases
  	 
Wide Column Stores	Key-Value Databases
   	    

@cloudtbt <http://www.anyannava.com>

# Agenda

Introduction

Why NoSQL?

Transactions, ACID properties and CAP theorem

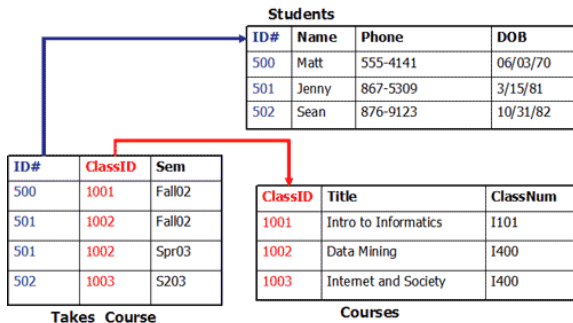
Data models

NoSQL databases design and implementation

# Relational databases

## SQL

- Born in the 70's – Still heavily used
- Data is organized into relations (in SQL: tables)
- Each relation is an unordered collection of tuples (rows)



# About SQL

## Advantages

- Separate the data from the code
  - ▶ High-level language
  - ▶ Space for optimization strategies
- Powerful query language
  - ▶ Clean semantics
  - ▶ Operations on sets
- Support for transactions

# Motivations for alternative models

see <https://blog.couchbase.com/nosql-adoption-survey-surprises/>

## Some limitations of relational databases

- Performance and scalability
  - ▶ Difficult to partition the data (in general run on a single server)
  - ▶ Need to scale up to improve performance
- Lack of flexibility
  - ▶ Will to easily change the schema
  - ▶ Need to express different relations
  - ▶ Not all data are well structured
- Few open source solutions
- Mismatch between the relational model and object-oriented programming model

# Illustration of the object-relational mismatch

Figure by M. Kleppmann

<http://www.linkedin.com/in/williamhgates>



**Bill Gates**  
Greater Seattle Area | Philanthropy

**Summary**  
Co-chair of the Bill & Melinda Gates Foundation. Chairman, Microsoft Corporation. Voracious reader. Avid traveler. Active blogger.

**Experience**  
Co-chair • Bill & Melinda Gates Foundation  
2000 – Present  
Co-founder, Chairman • Microsoft  
1975 – Present

**Education**  
Harvard University  
1973 – 1975  
Lakeside School, Seattle

**Contact Info**  
Blog: [thegatesnotes.com](http://thegatesnotes.com)  
Twitter: @BillGates

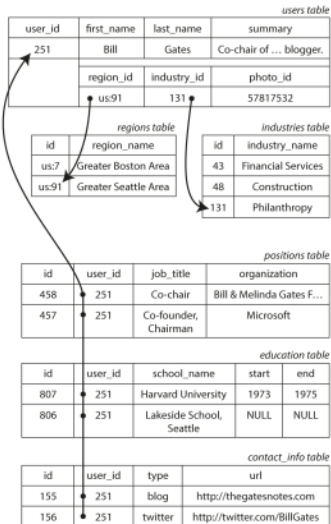


Figure: A CV in a relation database

# Illustration of the object-relational mismatch

Figure by M. Kleppmann

```
{
  "user_id": 251,
  "first_name": "Bill",
  "last_name": "Gates",
  "summary": "Co-chair of the Bill & Melinda Gates; Active blogger.",
  "region_id": "us:91",
  "industry_id": 131,
  "photo_url": "/p/7/000/253/05b/308dd6e.jpg",
  "positions": [
    {"job_title": "Co-chair", "organization": "Bill & Melinda Gates Foundation"},
    {"job_title": "Co-founder, Chairman", "organization": "Microsoft"}
  ],
  "education": [
    {"school_name": "Harvard University", "start": 1973, "end": 1975},
    {"school_name": "Lakeside School, Seattle", "start": null, "end": null}
  ],
  "contact_info": {
    "blog": "http://thegatesnotes.com",
    "twitter": "http://twitter.com/BillGates"
  }
}
```

Figure: A CV in a JSON document

# About NoSQL

## What is NoSQL?

- A hashtag
  - ▶ NoSQL approaches were existing before the name became famous
- No SQL
- New SQL
- Not only SQL
  - ▶ Relational databases will continue to exist alongside non-relational datastores



# About NoSQL

## A variety of NoSQL solutions

- Key-Value (KV) stores
- Wide column stores (Column family stores)
- Document databases
- Graph databases

## Difference with relational databases

There are several ways in which they differ from relational databases:

- Properties
- Data models
- Underlying architecture

# Agenda

Introduction

Why NoSQL?

Transactions, ACID properties and CAP theorem

Data models

NoSQL databases design and implementation

# About transactions

## The concept of transaction

- Groups several read and write operations into a logical unit
- A group of reads and writes are executed as one operation:
  - ▶ The entire transaction succeeds (commit)
  - ▶ or the entire transaction fails (abort, rollback)
- If a transaction fails, the application can safely retry

# About transactions

## The concept of transaction

- Groups several read and write operations into a logical unit
- A group of reads and writes are executed as one operation:
  - ▶ The entire transaction succeeds (commit)
  - ▶ or the entire transaction fails (abort, rollback)
- If a transaction fails, the application can safely retry

## Why do we need transactions?

- Crashes may occur at any time
  - ▶ On the database side
  - ▶ On the application side
  - ▶ The network might not be reliable
- Several clients may write to the database at the same time

# ACID

ACID describes the set of safety guarantees provided by transactions

- Atomicity
- Consistency
- Isolation
- Durability

Having such properties make the life of developers easy, but:

- ACID properties are not the same in all databases
  - ▶ It is not even the same in all SQL databases
- NoSQL solutions tend to provide weaker safety guarantees
  - ▶ To have better performance, scalability, etc.

# ACID: Atomicity

## Description

- A transactions succeeds completely or fails completely
  - ▶ If a single operation in a transaction fails, the whole transaction should fail
  - ▶ If a transaction fails, the database is left unchanged
- It should be able to deal with any faults *in the middle* of a transaction
- If a transaction fails, a client can safely retry

## In the NoSQL context:

- Atomicity is still ensured

# ACID: Consistency

## Description

- Ensures that the transaction brings the database from a valid state to another valid state
  - ▶ Example: Credits and debits over all accounts must always be balanced
- It is a property of the application, not of the database
  - ▶ The application cannot enforce application-specific invariants
  - ▶ The database can check some specific invariants
    - A foreign key must be valid

## In the NoSQL context:

- Consistency is (often) not discussed

# ACID: Durability

## Description

- Ensures that once a transaction has committed successfully, data will not be lost
  - ▶ Even if a server crashes (flush to a storage device, replication)

## In the NoSQL context:

- Durability is also ensured



# ACID: Isolation

## Description

- Concurrently executed transactions are isolated from each other
  - ▶ We need to deal with concurrent transactions that access the same data
- **Serializability**
  - ▶ High level of isolation where each transaction executes as if it was the only transaction applied on the database
    - As if the transactions are applied *serially*, one after the other
  - ▶ Many SQL solutions provide a lower level of isolation

In the NoSQL context:

- **What about the CAP theorem?**

# The CAP theorem

## 3 properties of databases

- Consistency
  - ▶ What guarantees do we have on the value returned by a read operation?
  - ▶ It strongly relates to Isolation in ACID (and not to consistency)
- Availability
  - ▶ The system should always accept updates
- Partition tolerance
  - ▶ The system should be able to deal with a partitioning of the network

## Comments on CAP theorem

- Was introduced by E. Brewer in its lectures (beginning of years 2000)
- Goal: discussing trade-offs in database design

# What does the CAP theorem says?


## The theorem

It is impossible to have a system that provides Consistency, Availability, and partition tolerance.

## How it should be understood:

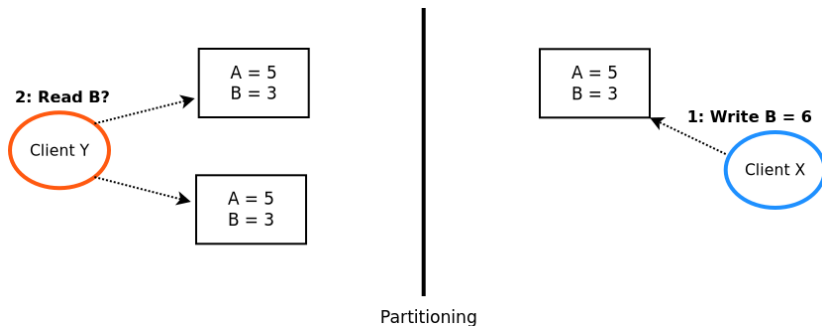
- Partitions are unavoidable
  - ▶ It is a fault, we have no control on it
- We need to choose between availability and consistency
  - ▶ In the CAP theorem:
    - Consistency is meant as *linearizability* (the strongest consistency guarantee)
    - Availability is meant as *total availability*
  - ▶ In practice, different trade-offs can be provided

# The intuition behind CAP



- Let inconsistencies occur? (No C)
- Stop executing transactions? (No A)

# The intuition behind CAP



- Let inconsistencies occur? (No C)
- Stop executing transactions? (No A)

Note that in a centralized system (non-partitioned relational database), no need for Partition tolerance

- We can have Consistency and Availability

# The impact of CAP on ACID for NoSQL

source: E. Brewer

## The main consequence

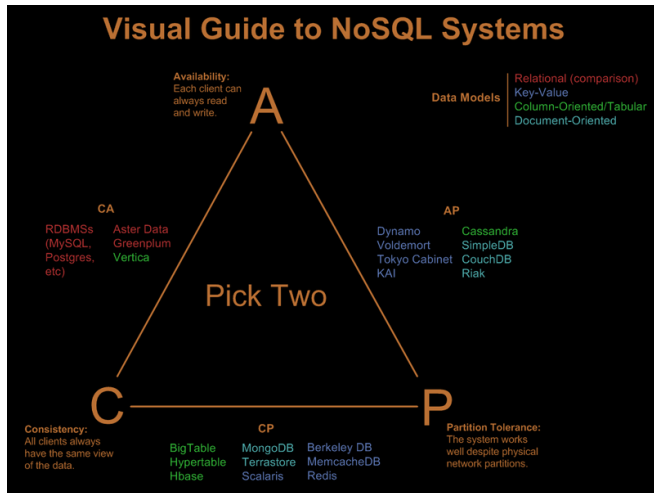
- No NoSQL database with strong Isolation

## Discussion about other ACID properties

- Atomicity
  - ▶ Each side should ensure atomicity
- Durability
  - ▶ Should never be compromised

# A vision of the NoSQL landscape

Source: <https://blog.nahurst.com/visual-guide-to-nosql-systems>



To be read with care:

- Solutions often provide a trade-off between CP and AP
- A single solution may often a different trade-off depending on how is is configured.
- **We don't pick two !**

# Agenda

Introduction

Why NoSQL?

Transactions, ACID properties and CAP theorem

Data models

NoSQL databases design and implementation




# Key-Value store

- Data are stored as key-value pairs
  - ▶ The value can be a data structure (eg, a list)
- In general, only support single-object transactions
  - ▶ In this case, key-value pairs
- Examples:
  - ▶ Redis
  - ▶ Voldemort
- Use case:
  - ▶ Scalable cache for data
  - ▶ Note that some solutions ensure durability by writing data to disk

# Key-value store

Image by J. Stolfi



# Column family stores

- Data are organized in rows and columns (Tabular data store)
  - ▶ The data are arranged based on the rows
  - ▶ Column families are defined by users to improve performance
    - Group related columns together
- Only support single-object transactions
  - ▶ In this case, a row
- Examples:
  - ▶ BigTable/HBase
  - ▶ Cassandra
- Use case:
  - ▶ Data with some structure with the goal of achieving scalability and high throughput
  - ▶ Provide more complex lookup operations than KV stores

# Column family stores

Order Table

RowKey 127698	<b>Family: Customer</b> FirstName Adam Surname Fowler MemberID 831642 Status Premier	<b>Family: Items</b> Item-4 2 Item-43 6 Item-9 1	<b>Family: Delivery</b> Notes Leave with Neighbor ETA 2014-12-23 09:00
RowKey 895482	<b>Family: Customer</b> FirstName Joe Surname Bloggs	<b>Family: Items</b> Item-32 1 Item-72 2	<b>Family: Delivery</b> ETA 2015-01-03 14:00
⋮			

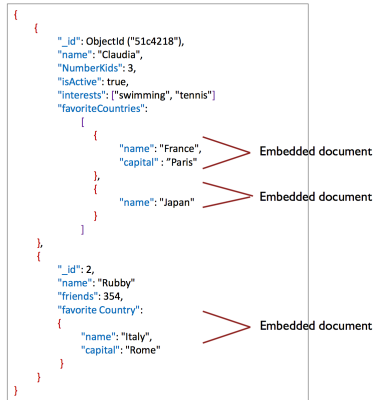
Note that not a row does not need to have an entry for all columns

# Document databases

- Data are organized in Key-Document pairs
  - ▶ A document is a nested structure with embedded metadata
  - ▶ No definition of a global schema
  - ▶ Popular formats: XML, Json
- Only support single-object transactions
  - ▶ In this case, a document or a field inside a document
- Examples:
  - ▶ MongoDB
  - ▶ CouchDB
- Use case:
  - ▶ An alternative to relational databases for structured data
  - ▶ Offer a richer set of operations compared to KV stores: Update, Find, etc.

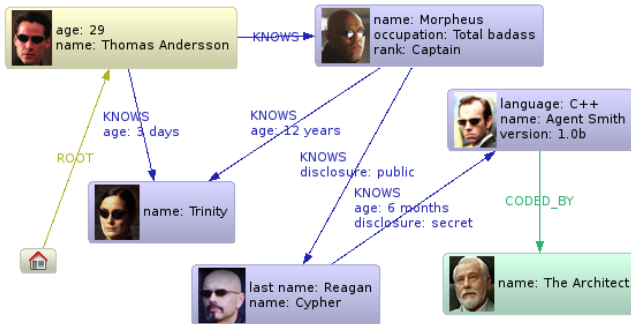
# Document DB

A document can have one or more documents inside.



# Graph DB


- Represent data as graphs
  - Nodes / relationships with properties as K/V pairs



# Graph DB: Neo4j

- Rich data format
  - Query language as pattern matching
  - Limited scalability
    - Replication to scale reads, writes need to be done to every replica

Cypher Query Language





# On the Many-to-one relationship

## Many-to-one relationship

- Many items may refer to the same item
- Example: Many people went to the same university

## Relational vs Document DB

# On the Many-to-one relationship

## Many-to-one relationship

- Many items may refer to the same item
- Example: Many people went to the same university

## Relational vs Document DB

- Relational databases use a foreign key
  - ▶ Consistency and low memory footprint (normalization)
  - ▶ Easy updates and support for joins
  - ▶ Difficult to scale
- Document databases duplicate data
  - ▶ Efficient read operations
  - ▶ Easy to scale
  - ▶ Higher memory footprint and updates are more difficult (risk of consistency issues)
    - Transactions on multiple objects could be very useful in this case
  - ▶ Join operations have to be implement by the application

# More on relations

## One-to-many relationship

- An item may have several entries of the same kind
- Example: One person may have had several positions during her career.
- Document DB allow storing such information easily and allow simple read operations

## Many-to-many relationship

- An item may have several entries of the same kind that are referred by multiple items
- Example: Several persons may have worked in the same company.
- Document DB may not have good support for such relationships

# Agenda

Introduction

Why NoSQL?

Transactions, ACID properties and CAP theorem

Data models

NoSQL databases design and implementation

# Google BigTable

- Column family data store
  - ▶ Data storage system used by many Google services: Youtube, Google maps, Gmail, etc.
- Paper published by Google in 2006 (F. Chang et al)
  - ▶ Now available as a service on Google Cloud
- Many ideas reused in other NoSQL databases



# Motivations

- A system that can store very large amount of data
  - ▶ TB or PB of data
  - ▶ A very large number of entries
  - ▶ Small entries (each entry is an array of bytes)
- A simple data model
  - ▶ Key-value pairs (A key identifies a row)
  - ▶ Multi-dimensional data
  - ▶ Sparse data
  - ▶ Data are associated with timestamps
- Works at very large scale
  - ▶ Thousands of machines
  - ▶ Millions of users

# About the data model

- Rows are identified by keys (arbitrary strings)
  - ▶ Modifications on one row are atomic
  - ▶ Rows are maintained in lexicographic order
- Columns are grouped in columns families
  - ▶ Columns can be sparsed
  - ▶ Clients can ask to retrieve a column family for one row
- Each cell can contain multiple versions indexed by a timestamp
  - ▶ Assigned by BigTable or by the client
  - ▶ Most recent versions are accessed first
  - ▶ GC politics:
    - Keep last n versions
    - Keep all new-enough versions

# About the data model

The diagram illustrates a data model table. On the left, a vertical arrow labeled "Sorted rows" points downwards. The table has four rows, each with a row key. The columns are grouped into four column families, indicated by brackets above the table: "language:", "contents:", "anchor:cnnsi.com", and "anchor:mylook.ca". The first column family contains the values "EN" for all rows. The second column family contains HTML doctype strings. The third column family contains the value "CNN" for the row with key "com.cnn.www". The fourth column family contains the value "CNN.com" for the row with key "com.cnn.www".

<i>row keys</i>	<i>column family</i> "language:"	<i>column family</i> "contents:"	<i>column family</i> anchor:cnnsi.com	<i>column family</i> anchor:mylook.ca
com.aaa	EN	<!DOCTYPE html PUBLIC...		
com.cnn.www	EN	<!DOCTYPE HTML PUBLIC...	"CNN"	"CNN.com"
com.cnn.www/TECH	EN	<!DOCTYPE HTML>...		
com.weather	EN	<!DOCTYPE HTML>...		



# Partitioning and performance

see <https://cloud.google.com/bigtable/docs/schema-design>

## Partitioning

- Partitioning on the rows
- Rows with close keys are in the same partition

## Recommendations about the schema for performance

- Accesses can be made based on key, key-prefix or key-range
  - ▶ Choose keys appropriately to make sequential accesses to a single host
  - ▶ Example: Reverse domain name, timestamps
  - ▶ To avoid: Domain name, hash values
  - ▶ Take advantage of the concept of key prefix
- Group related columns in a column family
  - ▶ Avoids retrieving all data from a single row when not needed
- Creating plenty of tables is not a good pattern
  - ▶ Use column families instead

# Sparse columns

see <https://cloud.google.com/bigtable/docs/overview>

"follows" column family

	Follows			
Row Key	gwasington	jadams	tjefferson	wmckinley
gwasington		1		
jadams	1		1	
tjefferson	1	1		1
wmckinley			1	

Multiple versions

# Building blocks of BigTable

## A master

- Assign tablets to servers
- With the help of a locking service


## Tablet servers

- Store the tables (divided in tablets)
- Process client requests

## Tablets

- Stored as SSTables (Sorted string tables)
- Stored in the Google File System for durability

# Implementation of tablets



# Implementation of tablets

## Write operations

- Data stored in memory ([Memtable](#))
- Any update is written to a commit log on GFS for durability
  - ▶ The log is shared between all hosted tablets

## Periodic writes to disk

- When the memtable becomes too big:
  - ▶ Copied as a new SSTable to GFS
    - Multiple SSTables are created if locality groups are defined (based on column families)
  - ▶ Reduces the memory footprint and reduces the amount of work to do during recovery
  - ▶ SSTables are immutable (no problem of concurrency control)
- Operation called [minor compaction](#)

# Implementation of tablets

## Read operations

- The state of the tablet = the memtable + all SSTables
  - ▶ A merged view needs to be created
  - ▶ The memtable and the SSTables may contain delete operations
- Locality groups help improving the performance of read operations

## Major compaction

- When the number of SSTables becomes too big, merge them into a single SSTable
  - ▶ Allow reclaiming resources for deleted data
  - ▶ Improve the performance of read operations

# Improving the performance of read operations

- During a read operation, potentially several SSTables need to be read
- How to avoid reading all SSTables when not needed?
  - ▶ Use of [Bloom filters](#)
  - ▶ Data structure that allows us to know if a SStable contains an entry for a given key-column pair

# Improving the performance of read operations

- During a read operation, potentially several SSTables need to be read
- How to avoid reading all SSTables when not needed?
  - ▶ Use of **Bloom filters**
  - ▶ Data structure that allows us to know if a SStable contains an entry for a given key-column pair


## Bloom filter

- Implements a membership function (is X in the set?)
- If the bloom filter answers no: it is guaranteed that X is not present
- If the bloom filter answers yes: the element is in the set with a high probability
- Good trade-off between accuracy and memory footprint



# About bloom filters

- A vector of  $n$  bits and  $k$  hash functions
- On insert:
  - ▶ Compute the  $k$  hash values
  - ▶ Set the corresponding bits to 1 in the vector
- On lookup:
  - ▶ Compute the  $k$  hash values
  - ▶ Test whether all bits are set to 1



## About the logs

On one node, a single commit log is created even if it hosts multiple tablets.

### Advantages

### Drawbacks

# About the logs

On one node, a single commit log is created even if it hosts multiple tablets.

## Advantages

- Write a single append-only file on disk
  - ▶ Improves performance by avoiding long seeks

## Drawbacks

- Recovery is more complex since the log includes data associated with different tablets
- The tablets might be distributed over multiple nodes

# Apache Cassandra

- Column family data store
- Paper published by Facebook in 2010 (A. Lakshman and P. Malik)
  - ▶ Used for implementing search functionalities
  - ▶ Released as open source
- Build on top of several ideas introduced by BigTable
  - ▶ **Warning:** Many changes in the design have been made since the first version of Cassandra



To be continued

# Additional references

## Mandatory reading

- *Bigtable: A Distributed Storage System for Structured Data.*, F. Chang et al., OSDI, 2006.
- *Cassandra: a decentralized structured storage system* ., A. Lakshman et al., SIGOPS OS review, 2010.

## Suggested reading

- <http://martin.kleppmann.com/2015/05/11/please-stop-calling-databases-cp-or-ap.html>, M. Kleppmann, 2015.
- <https://jvns.ca/blog/2016/11/19/a-critique-of-the-cap-theorem/>, J. Evans, 2016.