

Data Management in Large-Scale Distributed Systems

NoSQL Databases Fundamentals

Thomas Ropars

`thomas.ropars@univ-grenoble-alpes.fr`

`http://tropars.github.io/`

2020

References

- The lecture notes of V. Leroy
- The lecture notes of F. Zanen Boito
- Designing Data-Intensive Applications by Martin Kleppmann
 - ▶ Chapters 2 and 7

In this lecture

- Motivations for NoSQL databases
- ACID properties and CAP Theorem
- A landscape of NoSQL databases

Agenda

Introduction

Why NoSQL?

Transactions, ACID properties and CAP theorem

Data models

Common patterns of data accesses

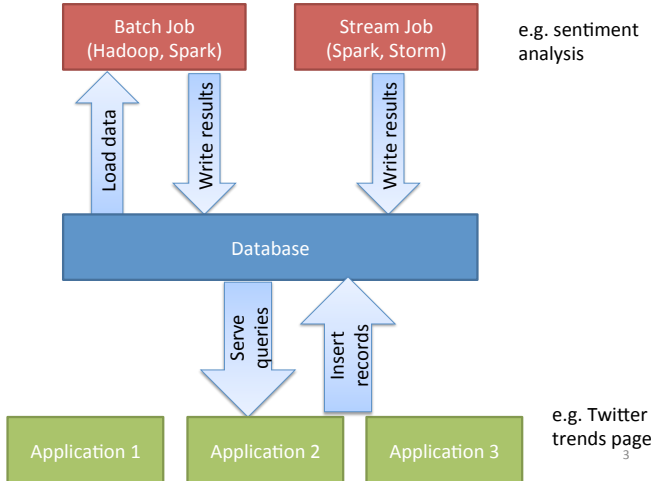
Large-scale data processing

- Batch processing: Hadoop, Spark, etc.
- Perform some computation/transformation over a full dataset
- Process all data

Selective query

- Access a specific part of the dataset
- Manipulate only the needed data
 - ▶ 1 record among millions
- Main purpose of a database system

Processing / Database Link



Different types of databases

- So far we used HDFS



- A file system can be seen as a very basic database
- Directories / files to organize data
- Very simple queries (file system path)
- Very good scalability, fault tolerance ...

- Other end of the spectrum: Relational Databases

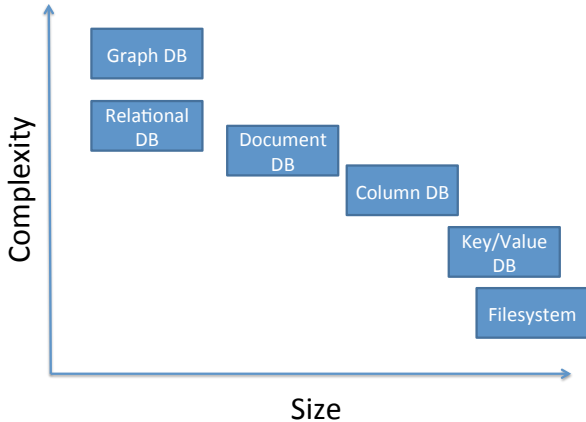


PostgreSQL















- SQL query language, very expressive
- Limited scalability (generally 1 server)



Size / Complexity



The NoSQL Jungle

Document Database	Graph Databases
  	 
Wide Column Stores	Key-Value Databases
   	    

@cloudtbt <http://www.anyannava.com>

Agenda

Introduction

Why NoSQL?

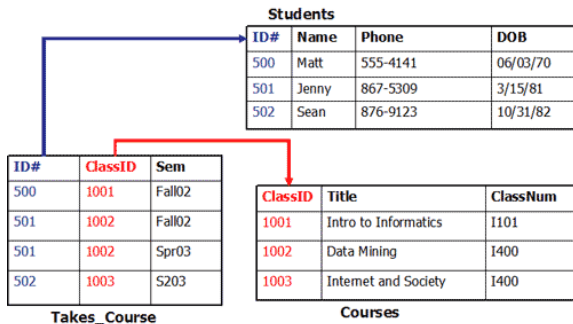
Transactions, ACID properties and CAP theorem

Data models

Relational databases

SQL

- Born in the 70's – Still heavily used
- Data is organized into relations (in SQL: tables)
- Each relation is an unordered collection of tuples (rows)
- Foreign keys are used to define the relationships among the tables



About SQL

Advantages

- Separate the data from the code
 - ▶ High-level language
 - ▶ Space for optimization strategies
- Powerful query language
 - ▶ Clean semantics
 - ▶ Operations on sets
- Support for transactions

Motivations for alternative models

see <https://blog.couchbase.com/nosql-adoption-survey-surprises/>

Some limitations of relational databases

- Performance and scalability
 - ▶ Difficult to partition the data (in general run on a single server)
 - ▶ Need to scale up to improve performance
- Lack of flexibility
 - ▶ Will to easily change the schema
 - ▶ Need to express different relations
 - ▶ Not all data are well structured
- Few open source solutions
- Mismatch between the relational model and object-oriented programming model

Illustration of the object-relational mismatch

Figure by M. Kleppmann

<http://www.linkedin.com/in/williamhgates>



Bill Gates
Greater Seattle Area | Philanthropy

Summary
Co-chair of the Bill & Melinda Gates Foundation. Chairman, Microsoft Corporation. Voracious reader. Avid traveler. Active blogger.

Experience
Co-chair • Bill & Melinda Gates Foundation
2000 – Present
Co-founder, Chairman • Microsoft
1975 – Present

Education
Harvard University
1973 – 1975
Lakeside School, Seattle

Contact Info
Blog: thegatesnotes.com
Twitter: @BillGates

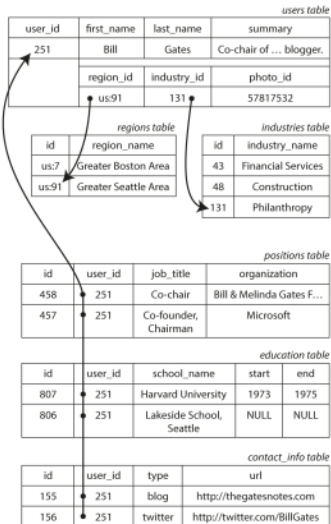


Figure: A CV in a relation database

Illustration of the object-relational mismatch

Figure by M. Kleppmann

```
{
  "user_id": 251,
  "first_name": "Bill",
  "last_name": "Gates",
  "summary": "Co-chair of the Bill & Melinda Gates; Active blogger.",
  "region_id": "us:91",
  "industry_id": 131,
  "photo_url": "/p/7/000/253/05b/308dd6e.jpg",
  "positions": [
    {"job_title": "Co-chair", "organization": "Bill & Melinda Gates Foundation"},
    {"job_title": "Co-founder, Chairman", "organization": "Microsoft"}
  ],
  "education": [
    {"school_name": "Harvard University", "start": 1973, "end": 1975},
    {"school_name": "Lakeside School, Seattle", "start": null, "end": null}
  ],
  "contact_info": {
    "blog": "http://thegatesnotes.com",
    "twitter": "http://twitter.com/BillGates"
  }
}
```

Figure: A CV in a JSON document

About NoSQL

What is NoSQL?

- A hashtag
 - ▶ NoSQL approaches were existing before the name became famous
- No SQL
- New SQL
- Not only SQL
 - ▶ Relational databases will continue to exist alongside non-relational datastores

About NoSQL

A variety of NoSQL solutions

- Key-Value (KV) stores
- Wide column stores (Column family stores)
- Document databases
- Graph databases

Difference with relational databases

There are several ways in which they differ from relational databases:

- Properties
- Data models
- Underlying architecture

Agenda

Introduction

Why NoSQL?

Transactions, ACID properties and CAP theorem

Data models

About transactions

The concept of transaction

- Groups several read and write operations into a logical unit
- A group of reads and writes are executed as one operation:
 - ▶ The entire transaction succeeds (commit)
 - ▶ or the entire transaction fails (abort, rollback)
- If a transaction fails, the application can safely retry

About transactions

The concept of transaction

- Groups several read and write operations into a logical unit
- A group of reads and writes are executed as one operation:
 - ▶ The entire transaction succeeds (commit)
 - ▶ or the entire transaction fails (abort, rollback)
- If a transaction fails, the application can safely retry

Why do we need transactions?

- Crashes may occur at any time
 - ▶ On the database side
 - ▶ On the application side
 - ▶ The network might not be reliable
- Several clients may write to the database at the same time

ACID

ACID describes the set of safety guarantees provided by transactions

- Atomicity
- Consistency
- Isolation
- Durability

Having such properties make the life of developers easy, but:

- ACID properties are not the same in all databases
 - ▶ It is not even the same in all SQL databases
- NoSQL solutions tend to provide weaker safety guarantees
 - ▶ To have better performance, scalability, etc.

ACID: Atomicity

Description

- A transactions succeeds completely or fails completely
 - ▶ If a single operation in a transaction fails, the whole transaction should fail
 - ▶ If a transaction fails, the database is left unchanged
- It should be able to deal with any faults *in the middle* of a transaction
- If a transaction fails, a client can safely retry

In the NoSQL context:

- Atomicity is still ensured

ACID: Consistency

Description

- Ensures that the transaction brings the database from a valid state to another valid state
 - ▶ Example: Credits and debits over all accounts must always be balanced
- It is a property of the application, not of the database
 - ▶ The database cannot enforce application-specific invariant
 - It is the responsibility of the programmer to issue transactions that make sense.
 - ▶ The database can check some specific invariant
 - A foreign key must be valid

In the NoSQL context:

- Consistency is (often) not discussed

ACID: Durability

Description

- Ensures that once a transaction has committed successfully, data will not be lost
 - ▶ Even if a server crashes (flush to a storage device, replication)

In the NoSQL context:

- Durability is also ensured

ACID: Isolation

Description

- Concurrently executed transactions are isolated from each other
 - ▶ We need to deal with concurrent transactions that access the same data
- **Serializability**
 - ▶ High level of isolation where each transaction executes as if it was the only transaction applied on the database
 - As if the transactions are applied *serially*, one after the other
 - ▶ Many SQL solutions provide a lower level of isolation
 - Example: Read committed (see next slides)

In the NoSQL context:

- **What about the CAP theorem?**

More about isolation

Alternative (weak) isolation levels

Read Committed


- **No dirty reads**: A read only sees data that has been committed
- **No dirty writes**: A write only overwrites data that has been committed.

Many SQL databases implement this level of isolation.

Example of allowed execution with Read Committed

Transaction to be executed (Increment a counter)


```
Begin transaction  
Read count  
count = count + 1  
Write count  
End transaction
```



Example of allowed execution with Read Committed

Transaction to be executed (Increment a counter)

```
Begin transaction  
Read count  
count = count + 1  
Write count  
End transaction
```



No dirty writes does not prevent all inconsistencies.

Example of allowed execution with Read Committed

2 Transactions executed concurrently

On the accounts of Alice:

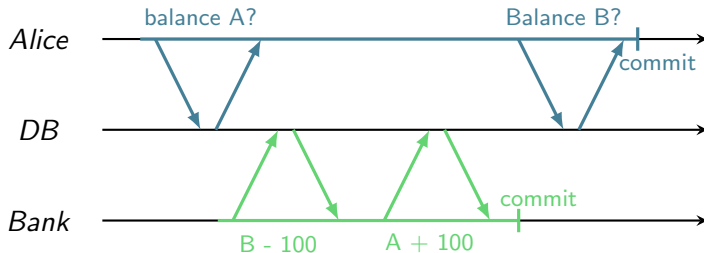
- initial state: \$500 of the two accounts
- Transaction 1: Transfer \$100 from account B to A
- Transaction 2: Check balance.

Example of allowed execution with Read Committed

2 Transactions executed concurrently

On the accounts of Alice:

- initial state: \$500 of the two accounts
- Transaction 1: Transfer \$100 from account B to A
- Transaction 2: Check balance.

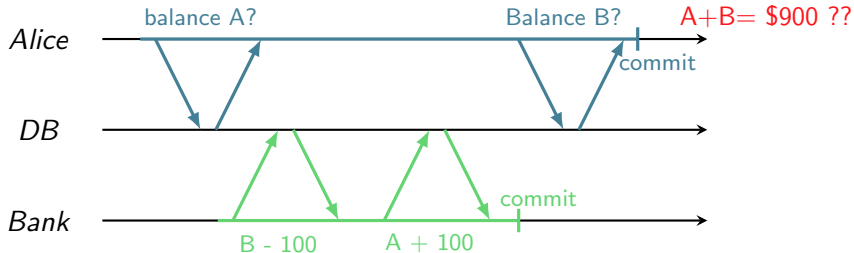


Example of allowed execution with Read Committed

2 Transactions executed concurrently

On the accounts of Alice:

- initial state: \$500 of the two accounts
- Transaction 1: Transfer \$100 from account B to A
- Transaction 2: Check balance.



Some inconsistencies can also be observed with [no dirty reads](#).

The CAP theorem

3 properties of databases

- Consistency
 - ▶ What guarantees do we have on the value returned by a read operation?
 - ▶ It strongly relates to Isolation in ACID (and not to consistency)
- Availability
 - ▶ The system should always accept updates
- Partition tolerance
 - ▶ The system should be able to deal with a partitioning of the network

Comments on CAP theorem

- Was introduced by E. Brewer in its lectures (beginning of years 2000)
- Goal: discussing trade-offs in database design

What does the CAP theorem says?

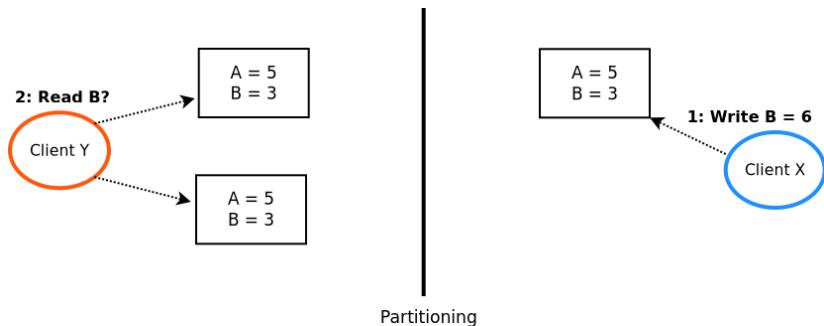
The theorem

It is impossible to have a system that provides Consistency, Availability, and partition tolerance.

How it should be understood:

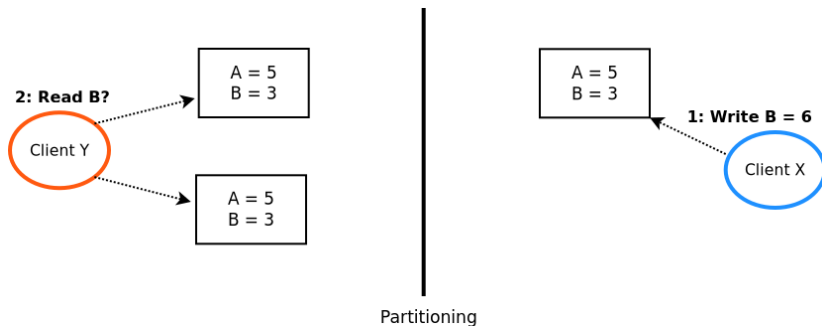
- Partitions are unavoidable
 - ▶ It is a fault, we have no control on it
- We need to choose between availability and consistency
 - ▶ In the CAP theorem:
 - Consistency is meant as *linearizability* (the strongest consistency guarantee)
 - Availability is meant as *total availability*
 - ▶ In practice, different trade-offs can be provided

The intuition behind CAP



- Let inconsistencies occur? (No C)
- Stop executing transactions? (No A)

The intuition behind CAP



- Let inconsistencies occur? (No C)
- Stop executing transactions? (No A)

Note that in a centralized system (non-partitioned relational database), no need for Partition tolerance

- We can have Consistency and Availability

The impact of CAP on ACID for NoSQL

source: E. Brewer

The main consequence

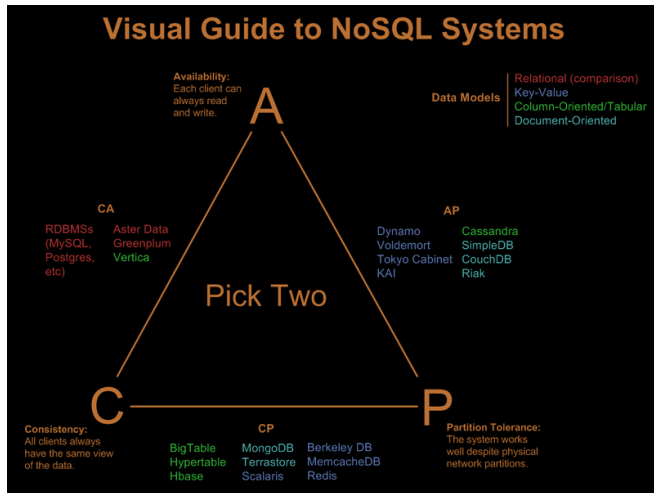
- No NoSQL database with strong Isolation

Discussion about other ACID properties

- Atomicity
 - ▶ Each side should ensure atomicity
- Durability
 - ▶ Should never be compromised

A vision of the NoSQL landscape

Source: <https://blog.nahurst.com/visual-guide-to-nosql-systems>



To be read with care:

- Solutions often provide a trade-off between CP and AP
- A single solution may offer a different trade-off depending on how is is configured.
- **We don't pick two !**

Agenda

Introduction

Why NoSQL?

Transactions, ACID properties and CAP theorem

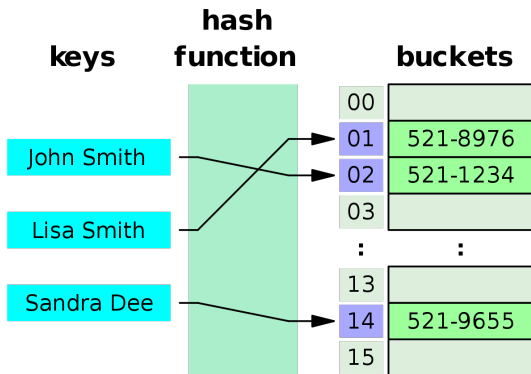
Data models

Key-Value store

- Data are stored as key-value pairs
 - ▶ The value can be a data structure (eg, a list)
- In general, only support single-object transactions
 - ▶ In this case, key-value pairs
- Examples:
 - ▶ Redis
 - ▶ Voldemort
- Use case:
 - ▶ Scalable cache for data
 - ▶ Note that some solutions ensure durability by writing data to disk

Key-value store

Image by J. Stolfi



Column family stores

- Data are organized in rows and columns (Tabular data store)
 - ▶ The data are arranged based on the rows
 - ▶ Column families are defined by users to improve performance
 - Group related columns together
- Only support single-object transactions
 - ▶ In this case, a row
- Examples:
 - ▶ BigTable/HBase
 - ▶ Cassandra
- Use case:
 - ▶ Data with some structure with the goal of achieving scalability and high throughput
 - ▶ Provide more complex lookup operations than KV stores

Column family stores

Order Table

RowKey 127698	Family: Customer FirstName Adam Surname Fowler MemberID 831642 Status Premier	Family: Items Item-4 2 Item-9 1 Item-43 6	Family: Delivery Notes Leave with Neighbor ETA 2014-12-23 09:00
RowKey 895482	Family: Customer FirstName Joe Surname Bloggs	Family: Items Item-72 2 Item-32 1	Family: Delivery ETA 2015-01-03 14:00
⋮			

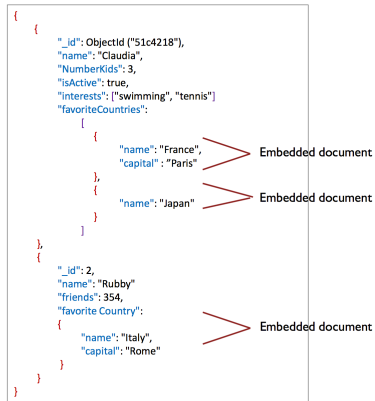
Note that not a row does not need to have an entry for all columns

Document databases

- Data are organized in Key-Document pairs
 - ▶ A document is a nested structure with embedded metadata
 - ▶ No definition of a global schema
 - ▶ Popular formats: XML, JSON
- Only support single-object transactions
 - ▶ In this case, a document or a field inside a document
- Examples:
 - ▶ MongoDB
 - ▶ CouchDB
- Use case:
 - ▶ An alternative to relational databases for structured data
 - ▶ Offer a richer set of operations compared to KV stores: Update, Find, etc.

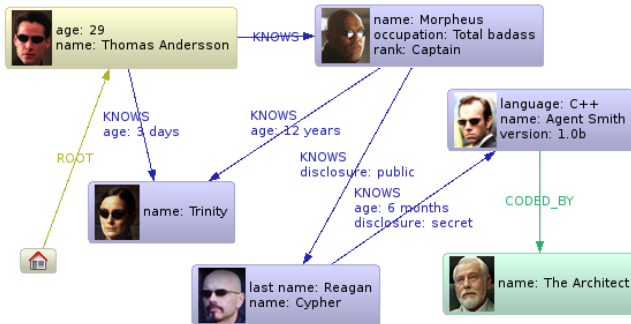
Document DB

A document can have one or more documents inside.



Graph DB


- Represent data as graphs
 - Nodes / relationships with properties as K/V pairs



Graph DB: Neo4j

- Rich data format
 - Query language as pattern matching
 - Limited scalability
 - Replication to scale reads, writes need to be done to every replica

Cypher Query Language



On the Many-to-one relationship

Many-to-one relationship

- Many items may refer to the same item
- Example: Many people went to the same university

Relational vs Document DB

On the Many-to-one relationship

Many-to-one relationship

- Many items may refer to the same item
- Example: Many people went to the same university

Relational vs Document DB

- Relational databases use a foreign key
 - ▶ Consistency and low memory footprint (normalization)
 - ▶ Easy updates and support for joins
 - ▶ Difficult to scale
- Document databases duplicate data
 - ▶ Efficient read operations
 - ▶ Easy to scale
 - ▶ Higher memory footprint and updates are more difficult (risk of consistency issues)
 - Transactions on multiple objects could be very useful in this case
 - ▶ Join operations have to be implemented by the application

More on relations

One-to-many relationship

- An item may have several entries of the same kind
- Example: One person may have had several positions during her career.
- Document DB allow storing such information easily and allow simple read operations

Many-to-many relationship

- An item may have several entries of the same kind that are referred by multiple items
- Example: Several persons may have worked in the same company.
- Document DB may not have good support for such relationships

Additional references

Suggested reading

- <http://martin.kleppmann.com/2015/05/11/please-stop-calling-databases-cp-or-ap.html>, M. Kleppmann, 2015.
- <https://jvns.ca/blog/2016/11/19/a-critique-of-the-cap-theorem/>, J. Evans, 2016.