

Data Management in Large-Scale Distributed Systems

Google BigTable

Thomas Ropars

`thomas.ropars@univ-grenoble-alpes.fr`

`http://tropars.github.io/`

2021

References

- The lecture notes of Prof. E. Brewer
- Designing Data-Intensive Applications by Martin Kleppmann
 - ▶ Chapters 2 and 7

In this lecture

- The design of Google BigTable
- Column Family Databases
 - ▶ Scalability, locality
- Solutions to efficiently store and retrieve data
 - ▶ SSTable (Sorted-String Table)
 - ▶ LSM Tree

Agenda

Introduction

Data Model

Storing Data

Distributed execution

Google BigTable

- Column family data store
 - ▶ Data storage system used by many Google services: Youtube, Google maps, Gmail, etc.
- Paper published by Google in 2006 (F. Chang et al)
 - ▶ Now available as a service on Google Cloud
- Many ideas reused in other NoSQL databases



Motivations

- A system that can store very large amount of data
 - ▶ TB or PB of data
 - ▶ A very large number of entries
 - ▶ Small entries (each entry is an array of bytes)
- A simple data model
 - ▶ Key-value pairs (A key identifies a row)
 - ▶ Sparse data
 - ▶ Multi-dimensional data: Data are associated with timestamps
- Works at very large scale
 - ▶ Thousands of machines
 - ▶ Millions of users

Agenda

Introduction

Data Model

Storing Data

Distributed execution

About the data model

- Rows are identified by keys (arbitrary strings)
 - ▶ Modifications on one row are atomic
 - ▶ Rows are maintained in lexicographic order
- Columns are grouped in columns families
 - ▶ Columns can be sparse
 - ▶ Clients can ask to retrieve a column family for one row
- Each cell can contain multiple versions indexed by a timestamp
 - ▶ Assigned by BigTable or by the client
 - ▶ Most recent versions are accessed first
 - ▶ GC politics:
 - Keep last n versions
 - Keep all new-enough versions

About the data model: An example

row keys

column family

column family

column family

Sorted rows

com.aaa

com.cnn.www

com.cnn.www/TECH

com.weather

language:	contents:	anchor:cnnsi.com	anchor:mylook.ca
EN	<!DOCTYPE html PUBLIC...		
EN	<!DOCTYPE HTML PUBLIC...	"CNN"	"CNN.com"
EN	<!DOCTYPE HTML>...		
EN	<!DOCTYPE HTML>...		

Figure: Example of Table: the *WebTable*

About the data model: An example

The diagram illustrates the WebTable data model. It features a table with four rows and four columns. The rows are labeled on the left as 'Sorted rows' with a downward arrow, and the row keys are 'com.aaa', 'com.cnn.www', 'com.cnn.www/TECH', and 'com.weather'. The columns are grouped into four column families: 'language:', 'contents:', 'anchor:cnnsi.com', and 'anchor:mylook.ca'. The 'language:' family contains the value 'EN' for all rows. The 'contents:' family contains HTML snippets for each row. The 'anchor:cnnsi.com' family contains the value 'CNN' for the 'com.cnn.www' row. The 'anchor:mylook.ca' family contains the value 'CNN.com' for the 'com.cnn.www' row.

	row keys	column family "language:"	column family "contents:"	column family anchor:cnnsi.com	column family anchor:mylook.ca
Sorted rows ↓	com.aaa	EN	<!DOCTYPE html PUBLIC...		
	com.cnn.www	EN	<!DOCTYPE HTML PUBLIC...	"CNN"	"CNN.com"
	com.cnn.www/TECH	EN	<!DOCTYPE HTML>...		
	com.weather	EN	<!DOCTYPE HTML>...		

Figure: Example of Table: the *WebTable*

- The row name is a reverse URL
- Multiple column families:
 - ▶ The column family *contents* can contain multiple version with timestamps (corresponding to when the content was fetched)
 - ▶ The column family *anchor* is sparse
 - For each existing web page (column), store the text of the anchor that points to the considered web page (row)

Partitioning and performance

see <https://cloud.google.com/bigtable/docs/schema-design>

Partitioning

- Partitioning on the rows
- Each partition includes a range of keys

Recommendations about the schema for performance

- Accesses can be made based on key, key-prefix or key-range
 - ▶ Choose keys appropriately to make sequential accesses to a single host
 - ▶ Example: Reverse domain name, timestamps, etc.
 - ▶ To avoid: Domain name, hash values
 - ▶ Take advantage of the concept of key prefix
- Group related columns in a column family
 - ▶ Avoids retrieving all data from a single row when not needed
- Creating plenty of tables is not a good pattern
 - ▶ Use column families instead

Agenda

Introduction

Data Model

Storing Data

Distributed execution

About data storage

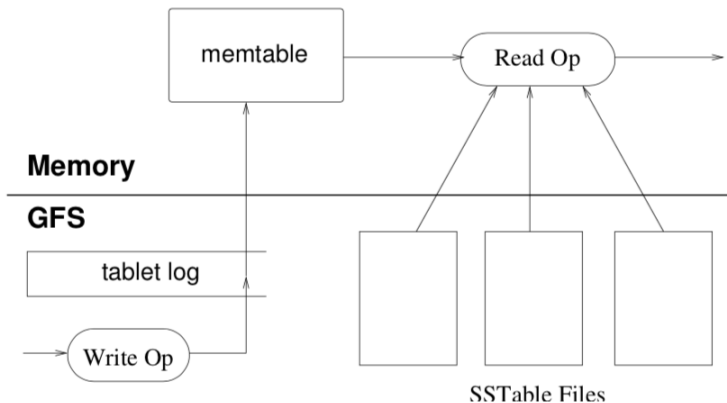
Main objectives and challenges

- Fast read and write operations
- Data persistence
- (Limit storage space consumption)

Tablets

- A partition of the table
 - ▶ Adjacent rows are stored in the same tablet
- Stored using SSTables
 - ▶ A persistent, ordered, immutable map
 - ▶ GFS is used as an underlying persistent file system
- Construction of an LSM (Log-Structured Merge) tree

Implementation of tablets



- **Tablet log**: Redo log of write operations
- **Memtable**: Sorted map of the most recent updates
- **SSTables**: Copy of previous memtables saved on Disk

SSTables

Sorted-String Table

Main principles

- File storing a sequence of key-value pairs (no padding)
 - ▶ Efficient sequential read/write operations
- An index can be created to know the offset of a key

Index

key	offset
key	offset
...	...

SSTable file

key	value	key	value	key	value
-----	-------	-----	-------	-----	-------	-----	-----

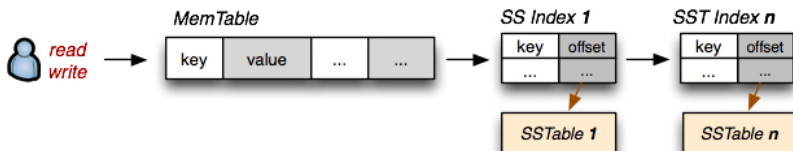
Figure: Figure by I. Grigorik

A Log-Structure Merge Tree

Figure by I. Grigorik

A multi-level data structure

- One fast level for new operations
 - ▶ In memory
 - ▶ The Memtable
- Other (slower) levels on disk
 - ▶ To save space in memory
 - ▶ Creation of a new level when the memory is full
 - Simple and efficient operation that flushes all content from memory to disk
 - Creation of an SSTable
 - Data on disk are immutable



Operation on Tablets

Write operations

- Data stored in memory ([Memtable](#))
- Any update is written to a commit log on GFS for persistence
 - ▶ The log is shared between all hosted tablets

Periodic writes to disk

- When the Memtable becomes too big:
 - ▶ Copied as a new SSTable to GFS
 - Multiple SSTables are created if locality groups are defined (based on column families)
 - ▶ Reduces the memory footprint and reduces the amount of work to do during recovery
 - ▶ SSTables are immutable (no problem of concurrency control)
- Operation called [minor compaction](#)

Implementation of tablets

Read operations

- The state of the tablet = the Memtable + all SSTables
 - ▶ A merged view needs to be created
 - The Memtable is accessed first, then the SSTables from most recent to oldest
 - ▶ The Memtable and the SSTables may contain delete operations
- Locality groups help improving the performance of read operations

Major compaction

- When the number of SSTables becomes too big, merge them into a single SSTable
 - ▶ Allow reclaiming resources for deleted data
 - ▶ Improve the performance of read operations

Improving the performance of read operations

- During a read operation, potentially several SSTables need to be read
 - ▶ How to avoid reading all SSTables when not needed?

An in-memory index

- Store information about the SSTables that contain information about a row
- Using a hash table
 - ▶ Constant time access
 - ▶ Potentially large memory footprint
- Using **Bloom filters**
 - ▶ Probabilistic data structure
 - ▶ Small memory footprint

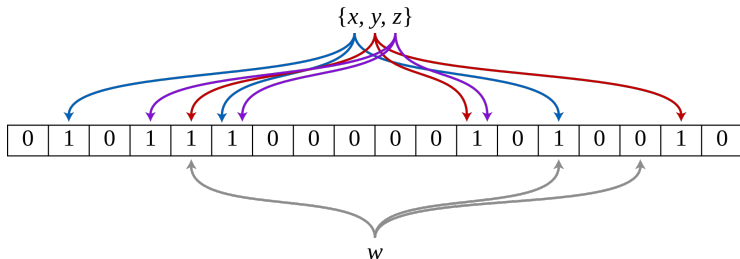
Improving the performance of read operations

Bloom filter

- Implements a membership function (is X in the set?)
- If the bloom filter answers no: it is guaranteed that X is not present
- If the bloom filter answers yes: the element is in the set with a high probability
- Good trade-off between accuracy and memory footprint

About bloom filters

- A vector of n bits and k hash functions
- On insert:
 - ▶ Compute the k hash values
 - ▶ Set the corresponding bits to 1 in the vector
- On lookup:
 - ▶ Compute the k hash values
 - ▶ Test whether all bits are set to 1



About the logs

On one node, a single commit log is created even if it hosts multiple tablets.

Advantages

Drawbacks

About the logs

On one node, a single commit log is created even if it hosts multiple tablets.

Advantages

- Write a single append-only file on disk
 - ▶ Improves performance by avoiding long seeks

Drawbacks

- Recovery is more complex since the log includes data associated with different tablets
- The tablets might be distributed over multiple nodes

Agenda

Introduction

Data Model

Storing Data

Distributed execution

Building blocks of the BigTable infrastructure

A master

- Assign tablets to servers
- With the help of a locking service

Tablet servers

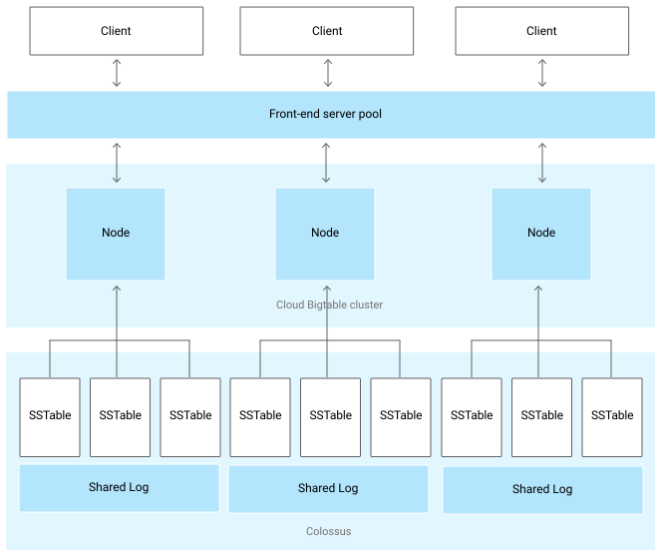
- Manage the tables (divided in tablets)
- Process client requests

Tablets

- Stores ranges of rows
- Stored as SSTables
- Stored in the Google File System for persistence

BigTable infrastructure

Image from <https://cloud.google.com/bigtable/docs/overview>



Additional references

Mandatory reading

- *Bigtable: A Distributed Storage System for Structured Data.*, F. Chang et al., OSDI, 2006.

Suggested reading

- <https://www.igvita.com/2012/02/06/ssstable-and-log-structured-storage-leveldb/>, I. Grigorik, 2012.