

# Data Management in Large-Scale Distributed Systems

File formats

Thomas Ropars

`thomas.ropars@univ-grenoble-alpes.fr`

`http://tropars.github.io/`

2020

# References

- Designing Data-Intensive Applications by Martin Kleppmann
  - ▶ Chapters 2: *Column-oriented storage*
  - ▶ Chapter 4: *Formats for Encoding Data*

## In this lecture

- Representation of large data on disks
  - ▶ Data that will be queried for analysis
- Column-oriented file formats

# Agenda

Introduction

Textual formats

Column-oriented formats

# Storing data on disks

- The representation of data on disks is in general not the same as in memory
  - ▶ Storing a pointer on disk would be meaningless
  - ▶ Random accesses on disk can be very slow
- Many file formats exist in the context of *Big Data*
  - ▶ CSV
  - ▶ JSON
  - ▶ Avro
  - ▶ Parquet
  - ▶ ORC
  - ▶ etc.

**What are the properties of each file format? Which one to choose?**

# Challenges

- Try to have a compact representation of the data
  - ▶ And organize the data so that they can be efficiently compressed
- Allow modifying the schema and ensure forward/backward compatibility
  - ▶ Not covered in this lecture
- Optimize the performance of read operations
  - ▶ Write once, read many

# Agenda

Introduction

Textual formats

Column-oriented formats

# Textual formats

## Examples of such formats

- CSV
- JSON
- XML

## Advantages

- Readable by humans

## Drawbacks

- High storage footprint
- Very low read performance



# Textual formats

## CSV

- Comma Separated Values
- Good for storing data organized as a single table
  - ▶ The name of the columns is given by the first row (not verbose)
- No hierarchical structure

## JSON – XML

- Support for hierarchical structures
- Very verbose (large footprint)

# Binary encoding formats

## Examples

- Avro (Hadoop)
- Thrift (Facebook)
- Protocol Buffers (Google)

## Idea

- Describe the data using a schema
- Pack all fields describing an item (a row) in a binary format

## Advantages

- Can lead to huge space reduction

# Example

By M. Kleppmann

```
{  
  "userName": "Martin",  
  "favoriteNumber": 1337,  
  "interests": ["daydreaming", "hacking"]  
}
```

Figure: A JSON document

## Storage space

- If stored as JSON text file: 81 bytes
- If stored as simple binary JSON encoding (not using a schema): 66 bytes
  - ▶ Space saved on the representation of numbers and on structure information

# With Avro

```
{
  "type": "record",
  "name": "Person",
  "fields": [
    { "name": "userName",
      "type": "string" },
    { "name": "favoriteNumber", "type": ["null", "long"], "default": null },
    { "name": "interests",
      "type": { "type": "array", "items": "string" } }
  ]
}
```

Figure: Definition of the schema

# With Avro

## Avro

Byte sequence (32 bytes):

0c	4d	61	72	74	69	6e	02	f2	14	04	16	64	61	79	64	72	65	61	6d
69	6e	67	0e	68	61	63	6b	69	6e	67	00								

Breakdown:

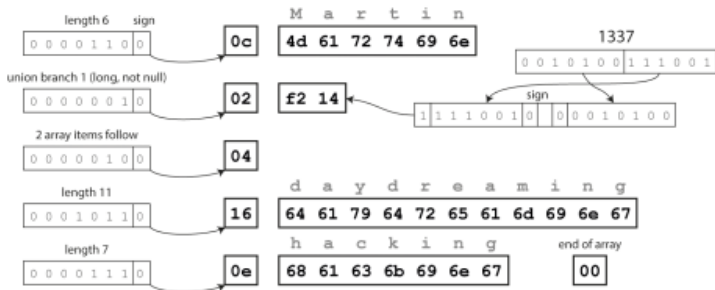


Figure: Binary representation of the item (32 bytes)

# Agenda

Introduction

Textual formats

Column-oriented formats

# Row-oriented formats

All formats described until now are row-oriented

- All the values from one row of a table are stored next to each other

## Limitations

# Row-oriented formats

All formats described until now are row-oriented

- All the values from one row of a table are stored next to each other

## Limitations

- Inefficient data compression
  - ▶ Data of different types are next to each other
- Inefficient read operations
  - ▶ We are often only interested in a few entries in a row
    - But we have to read the full row
  - ▶ We may want to filter elements based on a condition on one entry
    - But we have to read all the rows



# Column-oriented formats

## Examples

- Parquet (Twitter + Cloudera)
- ORC (Hadoop)

## Description

- Stores all the values from each column together
- Efficient compression
  - ▶ Close values are of the same type (e.g., integers)
  - ▶ The number of distinct values in a column is often small (not the case for rows)
- Optimizations on read:
  - ▶ Projection push-down
  - ▶ Predicate push-down

# Optimizations on read

## Projection push-down

- We are interested in a subset of columns
- We can read only the files corresponding to these columns
  - ▶ Or the file chunks storing these columns

## Predicate push-down

- We are interested in items corresponding to a condition
  - ▶ `SELECT * FROM Customers WHERE Country='Mexico'`
- Check the condition by reading only the corresponding column

# Parquet data layout

- Data are stored in **files**
- A file consists of one or more **row groups**
  - ▶ A set of rows
- A row group contains exactly one **column chunk** per column
  - ▶ A column chunk is contiguous in the file
- **Metadata** are stored at the end of the file
  - ▶ Position of each column chunk
  - ▶ Statistics about each chunk
    - Min/Max statistics for numbers
    - Dictionary filtering for other columns (as long as less than 40k different values)
- About **sorting**
  - ▶ Sorting rows based on the filtering criteria that is used the more often for filtering can improve performance

# Example

source: <https://blog.usejournal.com/sorting-and-parquet-3a382893cde5>

## Description of the data

- Customer table with one column being the country
- `SELECT * FROM Customers WHERE Country='Mexico'`
- Some numbers:
  - ▶ 10M rows
  - ▶ 10k rows per row group
  - ▶ 1% of the customers are from Mexico

## Amount of data read to answer the query

- With a row-based format:
- With an unsorted parquet file:
  
- With a sorted parquet file:

# Example

source: <https://blog.usejournal.com/sorting-and-parquet-3a382893cde5>

## Description of the data

- Customer table with one column being the country
- `SELECT * FROM Customers WHERE Country='Mexico'`
- Some numbers:
  - ▶ 10M rows
  - ▶ 10k rows per row group
  - ▶ 1% of the customers are from Mexico

## Amount of data read to answer the query

- With a row-based format: All data
- With an unsorted parquet file:
- With a sorted parquet file:

# Example

source: <https://blog.usejournal.com/sorting-and-parquet-3a382893cde5>

## Description of the data

- Customer table with one column being the country
- `SELECT * FROM Customers WHERE Country='Mexico'`
- Some numbers:
  - ▶ 10M rows
  - ▶ 10k rows per row group
  - ▶ 1% of the customers are from Mexico

## Amount of data read to answer the query

- With a row-based format: All data
- With an unsorted parquet file:
  - ▶ Probability of a row group with no customer from Mexico:  
 $0.99^{10000} = 2.25 \times 10^{-44}$
  - ▶ All row groups
- With a sorted parquet file:

# Example

source: <https://blog.usejournal.com/sorting-and-parquet-3a382893cde5>

## Description of the data

- Customer table with one column being the country
- `SELECT * FROM Customers WHERE Country='Mexico'`
- Some numbers:
  - ▶ 10M rows
  - ▶ 10k rows per row group
  - ▶ 1% of the customers are from Mexico

## Amount of data read to answer the query

- With a row-based format: All data
- With an unsorted parquet file:
  - ▶ Probability of a row group with no customer from Mexico:  
 $0.99^{10000} = 2.25 \times 10^{-44}$
  - ▶ All row groups
- With a sorted parquet file: 1% of the row groups (10)

# Another example

source: <https://www.linkedin.com/pulse/we-taking-only-half-advantage-columnar-file-format-eric-sun/>

## Description of the data

- A website log dataset
  - ▶ Information in one entry:  
timestamp, user\_id, cookie, page\_id, http\_header, ...
- Queries filters against page\_id

## Some results

- Avro:
  - ▶ Compressed data footprint: 1.4 TB
  - ▶ Amount of data read on query: 1.4 TB
- ORC<sup>1</sup> unsorted/sorted:
  - ▶ Compressed data footprint: 0.9 TB / 0.5 TB
  - ▶ Amount of data read on query: 300 GB / 200 MB

---

<sup>1</sup>similar to parquet



## Additional references

### Suggested reading

- *Dremel: Interactive Analysis of Web-Scale Datasets.*, S. Melnik et al., VLDB, 2010.