

Data Management in Large-Scale Distributed Systems

Storing Large Scale Graph Data

Thomas Ropars

`thomas.ropars@univ-grenoble-alpes.fr`

`http://tropars.github.io/`

2021

References

- The presentation slides of N. Bronson (TAO, Facebook's Distributed Data Store for the Social Graph)

In this lecture

- Managing Graph data
- Design of a large-scale geographically distributed database
 - ▶ Fast read requests
 - ▶ Low risks of inconsistency

Agenda

Introduction

Data Model

Working at scale

Facebook TAO

- Distributed Data Store for social graph
- Paper published by Facebook in 2013 (N. Bronson et al.)
 - ▶ Used to store and efficiently navigate through the data of a social media
 - ▶ A data model
 - ▶ An advanced replication + caching strategy to be able to go world scale
 - ▶ MySQL servers for storing data
- Evolution
 - ▶ Implement the database as a column-family LSM-tree based data store for better performance

Agenda

Introduction

Data Model

Working at scale

Social data representation

Representing people, actions and relationships

- = entities and connections
- Represented as nodes and edges in a graph

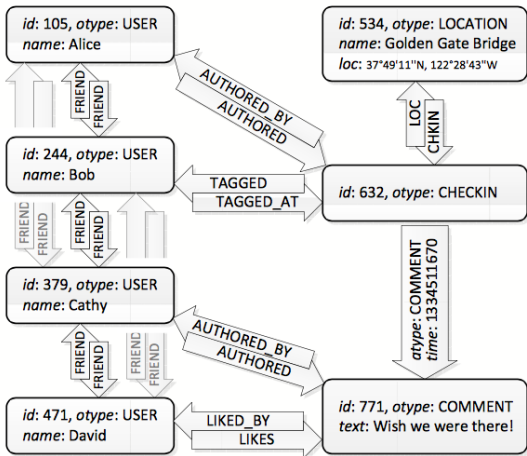


Data Model

Objects (e.g. user, place)
with unique IDs

Associations (e.g. tagged)
between two IDs

Both have key-value data
as well as a time field



API

Object/Association API

- Allocate, update, delete objects/associations

Association Query API

- Starting from a tuple: (object , associationType)
 - ▶ Association List: $(id1, atype) \rightarrow [a_{new} \dots a_{old}]$
 - ▶ Newer associations are returned first
- Examples of queries supported by the API:
 - ▶ `assoc_get(id1, atype, id2set, high?, low?)`
 - List associations between specific ids (with time bounds)
 - ▶ `assoc_range(id1, atype, pos, limit)`
 - Returns elements of the $(id1, atype)$ association list with $index \in [pos, pos + limit]$
 - The 50 most recent comment on Alice's checkin:
`assoc_range(632, COMMENT, 0, 50)`
 - ▶ `assoc_time_range(id1, atype, high, low, limit)`

Agenda

Introduction

Data Model

Working at scale


Challenges

- A very large dataset
- A large number of read and write requests
- Many geographically distributed data centers across the world.

Main objectives

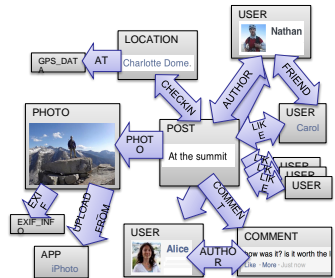

- Efficiency
 - ▶ Low read latency
- Consistency
 - ▶ Timeliness of writes
 - ▶ Achieve Read-After-Write consistency most of the time
- High read availability

Dynamically Rendering the Graph



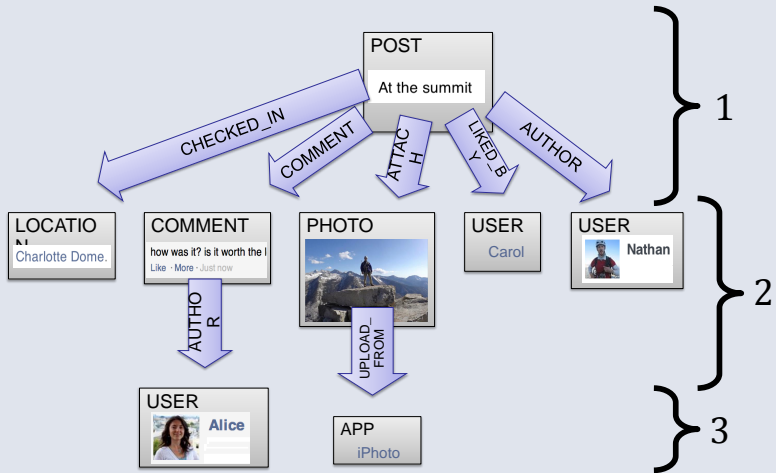
Web Server (PHP)

TAO



- 1 billion queries/second
- many petabytes of data

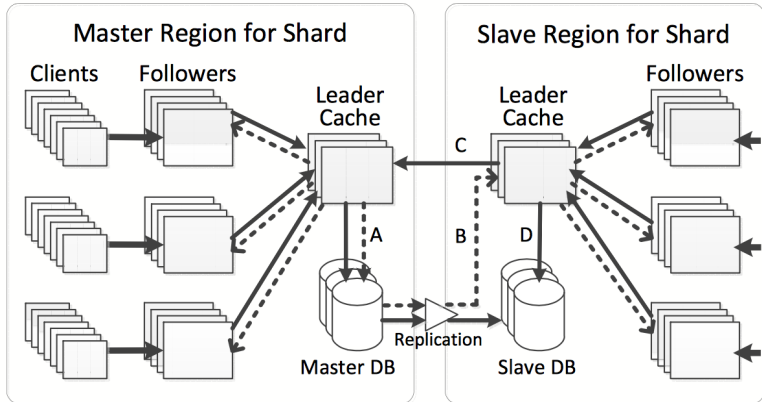
Dynamic Resolution of Data Dependencies



Design Principles

- The system is divided into regions
 - ▶ *Small* latency between data centers inside a region
 - ▶ A full copy of the social graph inside each region
- Replication of the data between regions
 - ▶ Master/slave replication
 - ▶ All writes go first to the master
- Data partitioning (sharding) over multiple database instances
- An in-memory cache is used to improve read performance
 - ▶ Based on Memcached KV store

Architecture



MySQL databases → durability

Leader cache → coordinates writes to each object

Follower caches → serve reads but not writes

TAO's caching architecture

Caching tiers

- Multiple servers
 - ▶ A set of servers form a caching tier
 - ▶ Data distributed based on sharding inside a tier
 - ▶ Clients send request to the correct server depending on the target object id.

A hierarchical architecture

- A single *leader* tier and multiple *follower* tiers
- A client contacts the closest follower tier

Read and write operations

Assuming a single region

Read requests

- Served by the follower caching tiers
- Forwarded to the master tier in case of miss

Write requests

- Forwarded to the master caching tier (write-through strategy)
 - ▶ Improves timeliness of writes
- The master caching tier orders the updates to the same objects and apply them to the database
 - ▶ The issuing follower is updated synchronously
 - ▶ The other followers are updated periodically (eventual consistency)

Write operations and geo-distribution

Some numbers

- Regions can be 1000's Kms away (high latency)
- 25 times more reads than writes

Geo-distribution

- Read requests are always served locally (inside a region)
 - ▶ Risk of stale data but low latency
- Write request always go to the master region
 - ▶ Requests forwarded to the leader caching tier
 - Simple replication protocol
 - Other databases replicas are updated asynchronously
 - ▶ Local leader cache updated synchronously

More on geo-distribution

Load balancing

- A different region has the leader role for different shards
- A region can be the leader for multiple shards

Locality

- An association is stored on the shard of its `id1`
- Association queries are served by a single server

Is consistency good enough?

- In practice, 99.99% of reads to vertices return a consistent result
- See "*Existential Consistency: Measuring and Understanding Consistency at Facebook*, 2015."

Additional references

Suggested reading

- *TAO: Facebook's Distributed Data Store for the Social Graph*, N. Bronson et al., USENIX ATC, 2013.