

Operating Systems

Redundant Array of Inexpensive Disks (RAID)

Thomas Ropars

`thomas.ropars@univ-grenoble-alpes.fr`

2017

References

The content of this lecture is inspired by:

- *Operating Systems: Three Easy Pieces* by R. Arpaci-Dusseau and A. Arpaci-Dusseau

Other references:

- *Modern Operating Systems* by A. Tanenbaum
- *Operating System Concepts* by A. Silberschatz et al.

Agenda

Introduction

RAID levels

Summary

Agenda

Introduction

RAID levels

Summary

Motivation

In a previous lecture, we have seen how disks basically work.

Motivation

In a previous lecture, we have seen how disks basically work.

Sometimes we would like more:

- Storage space
 - ▶ What if I have too many data for a single disk?
- Performance
 - ▶ What if I do a lot of reads and writes (I/O bound)?
- Reliability
 - ▶ What if my disk fails?

RAID

History

- A Case for Redundant Arrays of Inexpensive Disks (RAID) by D. Patterson, G. Gibson and R. Katz (1988)
 - ▶ Argue that RAID can perform better than expensive disks
 - ▶ Argue that this is true despite decreased MTTF (mean time to failure)
 - ▶ Defines 5 levels of RAID (still valid)
- Some manufacturers talk about “Redundant Array of Independent Disks”

RAID

Interface

- From the point of view of the OS, a RAID system is just a large, reliable, efficient disk.
- **Logical I/O**: The OS issues logical I/Os to the RAID system

Internals

- A standard connection (eg., SATA, SCSI)
- A set of disks
- Volatile memory for buffering
- Microcontroller(s) that operate the RAID logic
- **Physical I/O**: The RAID issues the physical I/Os to the disks

Metrics

- **Capacity**: Given N disks, how much client data can be stored?
- **Reliability**: How many disks failures may a design tolerate?
 - ▶ In the following, we assume a **fail-stop** failure model: a disk fails by crashing.
- **Performance**:
 - ▶ **Latency**: single-request latency
 - ▶ **Steady-state throughput**: total throughput of many concurrent requests (sequential or random)

Agenda

Introduction

RAID levels

Summary

RAID 0: Striping

- RAID level 0 is no RAID!
 - no redundancy
- Spread the blocks across the disks in round robin fashion.
 - With N disks, N blocks read/write in parallel
 - Often implemented at the granularity of blocks (other: bit-level, byte-level) – can be multiple blocks

Disk 0	Disk 1	Disk 2	Disk 3
0	2	4	6
1	3	5	7
8	10	12	14
9	11	13	15

Figure: Striping 16 blocks on 4 disks (2-block granularity)

RAID 0: Striping

Metrics

- Reliability: does not tolerate failures
- Capacity: $C \times N^1$
- Performance: optimal performance

Throughput²

Sequential Read: $N \times S$
Sequential Write: $N \times S$
Random Read: $N \times R$
Random Write: $N \times R$

Latency³

Read: T
Write: T

¹C: capacity of one disk; N: number of disks

²S/R: throughput of one disk with sequential/random accesses

³T: latency of read/write with a single disk

RAID 1: Mirroring

- RAID level 1 targets reliability
- It keeps several copies of each block
- Copies are stored on different disks

Disk 0	Disk 1	Disk 2	Disk 3
0	0	0	0
1	1	1	1
2	2	2	2
3	3	3	3

Figure: Mirroring 4 blocks on 4 disks

RAID 1: Mirroring

Metrics

- Reliability: tolerates $N - 1$ disk failures
- Capacity: C
- Performance:
 - ▶ With random read, multiple reads can be issued in parallel on different disks
 - ▶ Write latency increases because we have to wait for slowest disk to finish

Throughput

Sequential Read: S
Sequential Write: S
Random Read: $N \times R$
Random Write: R

Latency

Read: T
Write: $\geq T$

RAID 1+0: Mirroring + striping

- Apply striping across pairs of mirrored disks
- RAID 0+1 is also possible (mirror a striping array of disks)
 - ▶ Considered less reliable: one disk failure makes a full array unusable.

Disk 0	Disk 1	Disk 2	Disk 3
0	0	2	2
1	1	3	3
4	4	6	6
5	5	7	7

Figure: RAID 1+0 on 4 disks

RAID 1+0: Mirroring + striping

Metrics

- Reliability: tolerates 1 disk failure (worst case)
- Capacity: $N/2 \times C$
- Performance:
 - ▶ Throughput is improved thanks to striping
 - ▶ Mirrored disks can serve different requests on random reads
 - ▶ Write latency increases because we have to wait for slowest disk to finish (worst case seek and rotational delay)

Throughput

Sequential Read: $N/2 \times S$
Sequential Write: $N/2 \times S$
Random Read: $N \times R$
Random Write: $N/2 \times R$

Latency

Read: T
Write: $\geq T$

RAID 4: Parity-based redundancy

- Fault tolerance with reduced capacity lost
- Computes a parity block for each strip of blocks and store it on a separate disk
 - ▶ A bit-wise XOR is used to compute parity data (a parity block has the same size as a normal block)
 - ▶ If one disk fails, its data can be recovered based on the parity data and the data in the other disks

Disk 0	Disk 1	Disk 2	Disk 3	Disk 4
0	1	2	3	PP_0
4	5	6	7	PP_1
8	9	10	11	PP_2
12	13	14	15	PP_3

Figure: RAID level 4 on 5 disks

RAID 4: Parity-based redundancy

Metrics

- Reliability: tolerates 1 disk failure
- Capacity: $(N - 1) \times C$
- Performance:
 - ▶ For sequential writes, parity blocks can be written in parallel with the data stride
 - ▶ Writing a single block requires reading the block and the parity block first to be able to update the parity block
 - ▶ For random writes, the parity disk becomes the bottleneck (problem of small writes)

Throughput

Sequential Read: $(N - 1) \times S$
Sequential Write: $(N - 1) \times S$
Random Read: $(N - 1) \times R$
Random Write: $R/2$

Latency

Read: T
Write: $2T$

RAID 5: Rotating Parity

- Same advantages as RAID 4 but without the small writes performance issue
- Rotates the parity blocks across disks

Disk 0	Disk 1	Disk 2	Disk 3	Disk 4
0	1	2	3	PP_0
4	5	6	PP_1	7
8	9	PP_2	10	11
12	PP_3	13	14	15
PP_4	16	17	18	19

Figure: RAID level 5 on 5 disks

RAID 5: Rotating Parity

Metrics

- Reliability: tolerates 1 disk failure
- Capacity: $(N - 1) \times C$
- Performance:
 - ▶ Random reads can use all disks
 - ▶ Random writes allow using all disks in parallel but each write require 4 I/O operations (read/write of data and parity block)

Throughput

Sequential Read: $(N - 1) \times S$

Sequential Write: $(N - 1) \times S$

Random Read: $N \times R$

Random Write: $N/4 \times R$

Latency

Read: T

Write: $2T$

Agenda

Introduction

RAID levels

Summary

Summary

- Only interested in performance: RAID 0
- Random I/O performance and reliability: RAID 1+0
- Reliability and capacity: RAID 4
- Reliability and capacity + Random I/O performance: RAID 5

References for this lecture

- *Operating Systems: Three Easy Pieces* by R. Arpaci-Dusseau and A. Arpaci-Dusseau
 - ▶ Chapter 38: Redundant Disk Arrays (RAID)
- If you are interested in the topic:
 - ▶ D. Patterson, G. Gibson, and R. Katz. *A case for redundant arrays of inexpensive disks (RAID)*. ACM, 1988.