

# Parallel Algorithms and Programming

## Parallel algorithms in shared memory

Thomas Ropars

Email: [thomas.ropars@univ-grenoble-alpes.fr](mailto:thomas.ropars@univ-grenoble-alpes.fr)

Website: [tropars.github.io](http://tropars.github.io)

# References

The content of this lecture is inspired by:

- [Parallel algorithms](#) (Chapter 1) by H. Casanova, Y. Robert, A. Legrand.
- [A survey of parallel algorithms for shared-memory machines](#) by R. Karp, V. Ramachandran.
- [Parallel Algorithms](#) by G. Blelloch and B. Maggs.

# Outline

- The PRAM model
- Some shared-memory algorithms
- Analysis of PRAM models

# Need for a model

## A parallel algorithm


- Defines multiple operations to be executed in each step
- Includes communication/coordination between the processing units

## The problem

- A wide variety of parallel architectures
    - Different number of processing units
    - Multiple network topologies
- **How to reason about parallel algorithms?**
  - **How to avoid designing algorithms that would work only for one architecture?**
- A model can be used to abstract away some of the complexity
    - Should still capture enough details to predict with a reasonable accuracy how the algorithm will perform

# A model for shared memory computation

## The PRAM model



- *Parallel RAM*
- A shared central memory
- A set of processing units (PUs)
  - Any PU can access any memory location in one unit of time
- The number of PUs and the size of the memory is unbounded

# Details about the PRAM model

## Lock-step execution

- In each unit of time, a PU can:
  1. Read memory locations
  2. Run local computations
  3. Write to the shared memory
- All PUs execute these steps synchronously
  - No need for explicit synchronization

## About concurrent accesses to memory: 3 PRAM models

- **CREW**: Concurrent Read, Exclusive Write
- **CRCW**: Concurrent Read, Concurrent Write
  - Semantic of concurrent writes?
- **EREW**: Exclusive Read, Exclusive Write

# About the CRCW model

## Semantic of concurrent writes:

- *Arbitrary mode* : Select one value from the concurrent writes
- *Priority mode* : Select the value of the PU with the lowest index
- *Fusion mode* : A commutative and associative operation is applied to the values (logical OR, AND, sum, maximum, etc.)

## How powerful are the different models:

$$CRCW > CREW > EREW$$

A model is more powerful if there is one problem for which this model allows implementing a strictly faster solution with the same number of PUs

# Some shared-memory algorithms



# List ranking

## Description of the problem

- A linked list of  $n$  objects
  - Doubly-linked list
- We want to compute the distance of each element to the end of the list

## The sequential solution

- Iterate through the list from the end to the beginning
- Assign each element a distance from the last element while iterating

This solution has a complexity (execution time) in  $O(n)$

Can we do better with a parallel algorithm?

# List ranking

## A solution based on pointer jumping

```
# the list is stored in array *next*
# the distances are stored in array *d*
Ranking()
    forall i in parallel:                # initialization
        if next[i] is None:
            d[i] = 0
        else:
            d[i] = 1

    while there exists a node i such that next[i] != None:
        forall i in parallel do:
            if next[i] != None:
                d[i] = d[i] + d[next[i]]
                next[i] = next[next[i]] # pointer jumping
```

This solution has an execution time in  $O(\log n)$

- Note that the solution requires  $n$  PUs
- We note that the parallel version requires more work than the sequential version of the algorithm

# Comments on the previous algorithm

## Implementing pointer jumping

```
forall i in parallel:  
    next[i] = next[next[i]]
```

- In practice, if all processors do not execute synchronously, `next[next[i]]` may be overwritten by another PU before it is read here.
- To make the algorithm safe in practice, we would have to implement:

```
forall i in parallel:  
    temp[i] = next[next[i]]  
forall i in parallel:  
    next[i] = temp[i]
```

# Comments on the previous algorithm

## About the termination test

- Note that the test in the while loop can be done in constant time only in the CRCW model
- The problem is about having all PUs sharing the result of their local test (`next[i] != None`)
- In a **CW** model, all PUs can write to the same variable and a fusion operation can be used
- In a **EW** model, the results of the tests can only be aggregated two-by-two leading to a solution with a complexity in  $O(\log n)$  for this operation

# Point to root

## Description of the problem

- A tree data structure
- Each node should get a pointer to the root

## Use of pointer jumping

```
PointToRoot(P):  
    for k in 1..ceiling(log(sizeof(P))):  
        forall i in parallel:  
            P[i] = P[P[i]]
```

- We assume that we know sizeof(P)

# Scans (Prefix sums)

## Description of the problem

- Inputs:
  - A sequence of elements  $x_1, x_2 \dots x_n$
  - A associative operation  $*$
- Output:
  - A sequence of elements  $y_1, y_2 \dots y_n$  such that  $y_k = x_1 * x_2 \dots * x_k$

## Solution applying the pointer jumping technique

```
Scan(L):
    forall i in parallel:      # initialization
        y[i] = x[i]

    for k in 1..ceiling(log(sizeof(L))):
        forall i in parallel:
            if next[i] != None:
                y[next[i]] = y[i] * y[next[i]]
                next[i] = next[next[i]]
```

# Divide and conquer

- Split the problems into sub-problems that can be solved independently
- Merge the solutions

## Example: Mergesort

```
Mergesort(A):  
    if sizeof(A) is 1:  
        return A  
    else:  
        Do in parallel:  
            L = Mergesort(A[0 .. sizeof(A)/2])  
            R = Mergesort(A[sizeof(A)/2 .. sizeof(A)])  
        Merge(L,R)
```

It is usually important to parallelize the divide and the merge step:

- In the algorithm above, the merge step is going to be the bottleneck

# Analysis of PRAM models



# Comparison of PRAM models

## CRCW vs CREW

To compare CRCW and CREW, we consider a *reduce* operation over  $n$  elements with an associative operation.

- Example: the sum of  $n$  elements
  - With CRCW:  $O(1)$  steps
  - With CREW:  $O(\log n)$  steps

# Comparison of PRAM models

## CREW vs EREW

To compare CREW and EREW, we consider the problem of determining whether an element  $e$  belongs to a set  $(e_1, \dots, e_n)$ .

- Solution with CREW:
  - A boolean  $res$  is initialized to false and  $n$  PUs are used
  - PU  $k$  runs the test  $(e_k == e)$
  - If one PU finds  $e$ , it sets  $res$  to true
- Solution with EREW:
  - Same algorithm except  $e$  cannot be read simultaneously by multiple PUs
  - $n$  copies of  $e$  should be created (*broadcast*)

- With CREW:  $O(1)$  steps
- With EREW:  $O(\log n)$  steps

# Limits of the PRAM model

- Unrealistic memory model
  - Constant time access for all memory location
- Synchronous execution
  - Removes some flexibility